
Penn State University

Great Valley Campus

Engineering Division

Data Specification for

Data Warehouse – Part 1

Version 1.0

Date started - 09/15/2023

Table of Contents

Introduction.....	3
Purpose	3
Project Summary	3
Requirements Definition.....	4
Considerations.....	12
Document Change Log.....	12
Architecture Design.....	13
Entity Relationship Diagram	Error! Bookmark not defined.
Data Dictionary.....	13
Tables schemas.....	16
Examples of values	Error! Bookmark not defined.
ETL Process	21
ETL considerations	25
ETL Process Flow with description.....	25
Reporting System	32
Conclusions	44

Data Warehouse using PostgreSQL

by Eric Brown

INTRODUCTION

A data warehouse is a centralized repository for storing and managing large volumes of structured data from various sources within an organization. Thus, making it an integral part of an information system. Its primary purpose is to support business intelligence (BI) and data analytics activities by providing a structured and optimized environment for data storage, retrieval, and analysis. A data warehouse is typically made of several data sources, some of them remote. Though, data warehouses attempt to minimize their remote query selections.

While data warehouses bear many benefits from their overall data design infrastructure. One challenge data warehouses face happens during the design phase, which involves determining the correct views to materialize. Since, data warehouses have a very large, but finite storage capacity. So, it is not normally possible to store all data inside of them locally. This is why views included in data warehouses have to be chosen in manner that can answer queries of interest in a relatively fast time frame.

PURPOSE

Overall, the scope of the term project is to compare two different types of information systems. One being a data warehouse and one being a big data application framework called Hadoop. In this project we will focus on the creation and analysis of the data warehouse system using PostgreSQL.

PROJECT SUMMARY

This paragraph is used to introduce the following subsections, which can be used for an executive level overview.

A. Objectives

[Provide a concise description of the objectives of the document.]

Is to use the data from 2020 election contributions from the Propublica website to answer questions about the trends in election contributions amongst political parties.

B. Scope

[Briefly describe the scope of the requirements specification.]

The data requirements include the following:

- contributions must be for at least two candidates from each party
- contributions must be from at least four states of your choice
- contributions must be from three time periods which are to be loaded separately. These time periods can be weeks, months or quarters.

C. References

[Identify sources of information used to develop this document, such as IEEE or project documentation.]

- Project 1 Documentation
- Propublica
- PostgreSQL Documentation
- DBeaver Documentation
- KNIME Documentation

D. Outstanding Issues

[Provide a succinct list of issues or problems which are known to be outstanding with this revision.]

REQUIREMENTS DEFINITION

A. Goals

[[Provide a clear list of the expectations of the new application(s), both in terms of what must be improved and what must be retained from the current processes. All detailed requirements should address one or more of these goals]].

To provide a relational database management system (RDBMS) where the historical election data is stored in a star schema. And the necessary method to perform the ETL (Extract, Transform, Load) or CRUD (create, retrieve, update, delete) process of this system will be conducted through standard SQL commands. In conjunction to the use of Knime's as a tool to make the ETL process more efficient and less cumbersome. Knime has a wide variety of connectors and data transformation nodes to automate the ETL process.

Another goal/requirement for this project is to use the freeware open-source relational database management tool called PostgreSQL. This software will be used

because it is free to use, and its freeware nature makes it consistent with Hadoop. Which is a tool that is planned to be used in the next phase of the project. Also, PostgreSQL is highly extensible in terms of custom data types, functions, and operators. While also being ACID compliant in terms of data integrity and reliability. And being highly compatible across various types of operating systems. Whether it be used on Windows, macOS, or Linux, PostgreSQL runs in the same manner and provides the same security measures to protect data from unauthorized access across all operating systems.

A third goal of this project is to use PostgreSQL adherence to SQL standards to generate analysis on the historical data given in this project. This analysis would be conducted through different SQL commands and Knime's Report Writer tool ad-on to answer various prescribed business questions about the elections of 2020.

The last goal for this project is to use PostgreSQL as a freeware relational database management tool to stay consistent with the freeware nature of Hadoop. Mainly because the Hadoop software is to be used in the next portion of this project. Which involves using data warehouse systems to analyze big data.

B. Usability Requirements

[[Specify the requirements associated with ease of use, query capabilities, report layouts, online help and other interfaces to users and/or supervisors]].

Query Capabilities:

- PostgreSQL supports SQL (Structured Query Language) and provides a wide range of SQL features for querying and manipulating data. For example, users can perform complex queries on multiple tables and perform subqueries, joins, and aggregations on multiple tables
- PostgreSQL supports the creation and execution of stored procedures and user-defined functions using various procedural languages such as PL/pgSQL, PL/Python, and PL/Java
- PostgreSQL also allows for Window functions. Which is important because window functions do not cause rows to become grouped into a single output row — the rows retain their separate identities. Behind the scenes, the window functions can access more than just the current row of the query result. Which allows for advanced data analysis. As well as reporting by partitioning and ordering data within result sets.
- PostgreSQL also supports various indexing techniques, including B-tree, Hash, GIN, GiST, and SP-GiST, which can enhance a query's performance

Report Layouts:

- By itself PostgreSQL does not provide report layout capabilities. However, one could create custom reports by designing SQL queries to extract and format data in a manner that would be shown on a report
- Reporting tools like JasperReports, Tableau, or Power BI can be integrated with PostgreSQL to design and generate reports with advanced layout features
- You can use tools like LaTeX, in conjunction to PostgreSQL to generate formatted documents from database queries if you need complex report layouts
- In conjunction to PostgreSQL, Knime will be used mainly for construction the design for the report segment of this project
- Report Designer is an extension of the KNIME Analytics Platform. It is used to create custom reports and documents based on the data preprocessed through Knime.
- The Report Designer uses concepts of Business Intelligence and Reporting Tools (BIRT) to generate the overall structure of the report through two segments. The Master Page, which contains information for headers, footers, and titles common to every page in the report. And the Layout, which details how tables, charts, images, text, and other items are arranged.
- One could also bind report elements to one's KNIME data. Meaning the content of one's report could be dynamically updated as one's data or analysis changes. This is particularly useful when you have regularly updated data and want to automate the reporting process.

Online Help:

- PostgreSQL has extensive online documentation available on its official website. This documentation covers installation, configuration, SQL reference, and more
- One tool that could be used as help is PostgreSQL's pg_dump utility. It could be used to generate a SQL script containing the schema and data of a given database, which can serve as documentation to one's database
- The KNIME Analytics Platform has many online resources for troubleshooting, learning, and community engagement. Here are some of the common online support resources for KNIME: KNIME Forum, KNIME

Documentation, and KNIME Hub. As well as KNIME Blog and KNIME events/webinars

Interfaces to Users:

- PostgreSQL provides a command-line tool called psql that allows users to interact with the database using SQL commands through the command line
- There are several GUI tools available, such as pgAdmin and DBeaver that provide a user-friendly interface for interacting with a PostgreSQL database
- Developers can integrate PostgreSQL into applications using programming languages like Python, Java, and PHP by using appropriate libraries and drivers
- Web applications can be built using frameworks like Django, Ruby on Rails, or Node.js, with PostgreSQL as the backend database
- Users can connect to PostgreSQL databases through ODBC (Open Database Connectivity) or JDBC (Java Database Connectivity) drivers, enabling database interaction from various applications
- PostgreSQL has both web and mobile accessibility through REST API
- KNIME has various user interfaces for users to interact with the platform as well. These interfaces include KNIME's desktop interface, webportal, executor, and REST API

Additional Requirement to ease computer memory:

1.) On a Windows computer, if one sets docker with WSL, one needs to set some reasonable resource constraints on what WSL2 can actually use.

2.) Create a .wslconfig file (no extension) with the following:

- memory=4GB

IMPORTANT: no more than half of the ram in your computer, minimum 2.5 GB

- processors=2

IMPORTANT: no more than half the virtual processors you have

3.) To confirm these optimization changes, one should Open Windows PowerShell with admin rights, restart WSL2 by typing, *Restart-Service LxssManager*

C. System Security Requirements

[[Provide details of the security classification of the data handled by the system, special handling required for the data, and the types and levels of protection and control required for user access to the data]].

A Docker image runs a PostgreSQL database server by:

- (1) pulling the base of a Linux distribution such as Ubuntu
- (2) modifying it to include the PostgreSQL database and all other dependencies
- (3) making a new image out of that can be pushed to the Docker Hub

Once Docker creates the image, it can be used to create a container. Then the container is created in the user's file system and a read-write layer is added to the image. Also, the creation of a docker file can allow a network interface to be created to allow the local host to be connected through a user's port. This way an IP address is attached to the container and PostgreSQL database can run through the command line. In our project the local host will be through port 49168 on the user's local machine.

In terms of Third-Party Extensions:

Security measures will be in place based on the initial connection to DBEaver through the following authorization.

Host: localhost

Port: 49168

Database/Schema: student

User name: student

Password: student

D. Business Questions

[[Identify the business questions that should be answered by the software product. Include any prioritization of these requirements]].

- 1.) Which campaigning party had the highest contributions in terms of dollars between the four states of California, New Jersey, Florida, and Arizona
- 2.) Which campaigning party had the lowest contributions in terms of dollars amongst the four states of California, New Jersey, Florida, and Arizona
- 3.) Which state had the highest contributions to the top two campaigning parties
- 4.) Contributions made in California, New Jersey, Florida, and Arizona aggregated monthly in Q1 during the year of 2019

- 5.) Which company made the most contributions in terms of dollars to the top campaigning party
- 6.) What was the overall demographic makeup of the contributors of each of the top campaigning parties

E. Data Requirements

[[Identify the data elements and logical data groupings that will be required to answer the business questions above. Include any archiving requirements and the projected level of effort. This section may be supported by a data model and a data dictionary which should be included in an appendix]].

Typically, there are four steps in developing a data warehouse:

1. **Select the business process:** *This step is completed by converting the database schema into to a set of dimensional models for each discrete business processes.*

The main question for this project would be to provide information regarding contribution totals for each of the campaigning parties during different time periods of the 2020 election.

2. **Declare the grain:** *This step identifies exactly what each record in the fact table represents. For example, we may be interested in daily sales figures.*

For this project, we will choose monthly aggregation totals. This implies that we need to aggregate all our data sharing the same dimensions in the same month.

3. **Identify the dimensions:** *This step identifies those aspects of the facts that are interesting for the analytical process.*

filing_id = Normally, we would use this as a degenerate dimension but it is not aggregated sequentially, so it is irrelevant data

tran_id = We can use this as a degenerate dimension

linenumber = This a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

flag_orgind = This a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

date = This is a good dimension that define the grain of our data

cycle = Could be used a dimension to define time period, but it is a static data point

employer = This is a good dimension if we want to identify which employer are active in giving campaign contributions

occupation = This is a good dimension if we want to identify which occupation are active in giving campaign contributions

committee_name = This is an important dimension because it identifies which candidates a given contributor supports in the 2020 elections

org_name = This is an irrelevant dimension since it has no data list inside of it

address_one = This is a dimension that I will consider irrelevant in our case because we want to look only at trends and not at predicting individual behavior

address_two = This is a dimension that I will consider irrelevant in our case because we want to look only at trends and not at predicting individual behavior

last_name = This is a dimension that I will consider irrelevant in our case because we want to look only at trends and not at predicting individual behavior

first_name = This is a dimension that I will consider irrelevant in our case because we want to look only at trends and not at predicting individual behavior

middle_name = This is a dimension that I will consider irrelevant in our case because we want to look only at trends and not at predicting individual behavior

prefix = This is a dimension that could provide information regarding which gender contributed the most during the 2020 elections. But I would still consider it irrelevant for our business questions established previously. Also, there are many null values in this dimension

suffix = This is a dimension that could provide information regarding which profession contributed the most during the 2020 elections. But I would still consider it irrelevant for our business questions established previously. Mainly because most of the values in this dimension are null values

state = This is a good dimension if we want to identify which part of the country supports a particular type of candidate

zip = This is a good dimension if we want to delve deeper into identifying which part of the country supports a particular type of candidate

city = This is a good dimension if we want to delve deeper into identifying which part of the country supports a particular type of candidate

amount = This is not a dimension but rather a fact

aggregate_amount = This is not a dimension but rather a fact. But I would probably not want to use this aggregate total because it aggregates the amount field in a different manner than what we need for our business questions

memo_code = This is a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

memo_text = This is a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

back_ref_tran_id = This is a dimension that has similar information to another dimension in our database, tran_id. So, I would consider it not necessary to include as another dimension in our data warehouse

back_ref_sched_name = This a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

prigen = This a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

fecid = This a dimension that is irrelevant, since it has information that has no importance to our business objectives/findings

In conclusion we should create a star schema with eight dimensions:

- Tran_id
- Date
- Employer
- Occupation
- Committee_name
- State
- Zip
- City

For each of these dimensions we will create a table containing an id and a description field.

4. **Identify the facts:** *This step identifies what actionable data we need to store.*

Amount = is the main fact that we are trying to create our analysis around; total monetary contributions that were made to each political party

F. Design Constraints

[[Document any design constraints that should be taken into consideration during the system design phase]].

There must be at least three tables loaded into the PostgreSQL database. Create a star schema with eight dimensions:

- Tran_id
- Date
- Employer
- Occupation
- Committee_name
- State
- Zip
- City

CONSIDERATIONS

(May include as separate document)

DOCUMENT CHANGE LOG

Change Date	Version	CR #	Change Description	Author and Organization
09/24/23	1.0		Initial creation	Eric Brown
10/05/2023	2.0		Add more table details about data warehouse	Eric Brown
10/11/2023	3.0		Add schema details about data warehouse	Eric Brown
10/20/2023	4.0		Add information about ETL process	Eric Brown
10/14/2023	5.0		Add information regarding the report design and discoveries learned	Eric Brown
10/15/2023	6.0		Final edits for Project 1 submission	Eric Brown

2. ARCHITECTURE DESIGN

2.1 Relational Data Warehouse

Data Dictionary

A Data Dictionary is a document that describes the basic organization of a database. Typically a data dictionary will contain a list of variables in the database as well as the assigned variable names and a description of each type of variable. The data dictionary should also include the values accepted for each variable and any helpful comment such as important exclusions and skip patterns. The data dictionary is used primarily for data analysis.

The relational database was made up of four tables that had different contribution information inside of them. The four tables were based on the first quarter contribution results amongst four states. These states were Arizona, New Jersey, California, and Florida. All four tables were organized in the same way and listed below was generally how they were structured.

Contribution Table – Data Dictionary				
<i>Variable</i>	<i>Variable name</i>	<i>Variable type</i>	<i>Values</i>	<i>notes</i>
Filing Id	filing_id	String	SA17A, SA18, SB28A	Can not be null
Line number	linenum ber	String	SA17A, SA18, SB28A	Can not be null
Organization type	flag_organ d	String	IND	
Organization Name	org_name	String	ADP, 1144 Summit Ave, LLC, Atlantic Agri Imports, LLC	Has mostly null values
Contributor Last name	fname	String	Doe	Can not be null
Contributor First name	lname	String	John	Can not be null
Contributor middle name	middle_name	String	Bruce	
Professional abbreviations	prefix	String	Dr., Mrs, Mr, Judge	

Contribution Table – Data Dictionary				
Professional & Personal abbreviations	suffix	String	III, Jr., JR., M.D.	
Address	address_one	String	12 Larson Rd	
Second Part of Address	address_two	String	Unit 45, Suite 100	
City of Contributor	city	String	Phoenix, Los Angeles, Newark, Orlando	
State of Contributor	state	String	California, New Jersey, Florida, Arizona	must have a California, New Jersey, Arizona, or Florida state String value
Zipcode of Contributor	zip	String	85327	
Employer of Contributor	employer	String	Stagecoach Digital, City University Of New York- John Jay, etc..	
Occupation of Contributor	occupation	String	Engineer, Mechanic, Software programmer, etc...	
Dollar amount of contribution	amount	Numeric	2800, 1000, 26	
Date of contribution	date	Date/Time	3/6/2019, 3/5/2019, 3/31/2019, etc...	

Contribution Table – Data Dictionary				
Aggregated Amount contributed by contributor	aggregate_amount	Numeric	2800, 262.5 5000	
Memo code	memo_code	String	Null, X	Mostly null values
Memo text	memo text	String	* Earmarked Contribution through ACTBLUE on 01/13/2019, SEE REDESIGNATION	
Transaction Id of contribution	tran_id	Integer	819053, 815124, 815077	
Additional transaction Info	back_ref_transaction_id	String	1688351E, 169662E, 1691257E	
Campaign Contribution Year Id	prigen	String	P2020, G2020, Blank	
Campaign Year	cycle	Date	2020	Must be for year 2020
Id of Committee	fecid	String	C00698258 , C00694455	
Committee Name	committee_name	String	Kamala Harris For The People, Bernie 2020, Friends Of John Delaney	

Tables schemas

Provide a description of the physical schema of the data warehouse. Use the steps in the lesson and explain.

<i>Name of the table</i>	State		
Description	This table describes information regarding each state of the contributors.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique state	Integer	Between 1 and 12000
name	Name of state	String	California, New Jersey, Arizona, Florida
Primary Key	id		
Foreign Keys	NA		

<i>Name of the table</i>	City		
Description	This table describes information regarding each city of the contributors.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique city	Integer	Between 1 and 12000
city	Name of the city	String	Los Angeles, Key Largo, Tempe, Montclair
Primary Key	id		
Foreign Keys	NA		

<i>Name of the table</i>	Zipcode		
Description	This table describes information regarding each zipcode of the contributors.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique zipcode	Integer	Between 1 and 12000
zip	Zipcode of the city	Integer	33095, 7950, 85266, 90274
Primary Key	id		
Foreign Keys	NA		

<i>Name of the table</i>	Employer		
Description	This table describes information regarding each employer of the contributors.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique employer	Integer	Between 1 and 12000
employer	Employer name of each contributor	String	Town Of Prescott, Self Employed, Whole foods, Not Employed
Primary Key	id		
Foreign Keys	NA		

<i>Name of the table</i>	Occupation		
Description	This table describes information regarding each occupation of the contributors.		
Attribute	Description	Type	<i>Examples of values</i>

id	Id of each unique occupation	Integer	Between 1 and 12000
employer_title	Occupation title of each contributor	String	Sales, Planner, Attorney, Investor
Primary Key	id		
Foreign Keys	NA		

<i>Name of the table</i>	Date		
Description	This table displays all the contribution dates made by contributors.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique employer	Integer	Between 1 and 12000
date	Contribution transaction date	Date	9/30/2019, 10/1/2019, 11/4/2019
year	year portion of the date value		2019
quarter	quarter value of the year in which the date was in		1 to 4
month	month portion of the date value		1 to 12 1 = January 12 = December
weekofyear	week value of the date		ranging from 1 to 52
dayofweek	day value of the date value		ranging from 1 to 7 1 being equal to Monday & 2 being equal to Tuesday, etc..

Primary Key	id
Foreign Keys	NA

<i>Name of the table</i>	Transactions		
Description	This table displays all the contribution transactions made by contributors.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique contribution transaction	Integer	Between 1 and 12000
tran_id	Contribution transaction reference	Integer	1764024, 1964608, 454961
Primary Key	id		
Foreign Keys	NA		

<i>Name of the table</i>	Committee_Name		
Description	This table describes information regarding each committee in which the contributors made contributions.		
Attribute	Description	Type	<i>Examples of values</i>
id	Id of each unique committee name	Integer	Between 1 and 12000
committee_name	Names of the politicians that contributions were made to in the 2020 elections	String	Kamala Harris For The People, Bernie 2020, Friends Of John Delaney
Primary Key	id		
Foreign Keys	NA		

<i>Name of the Fact table</i>	AmountFacts			
Description	This table will store the fact variable - contribution amount in the data warehouse			
Attribute	Description	Type	<i>Examples of values</i>	<i>Notes</i>
stateid	State id	Integer	10	From the dimension table
cityid	City id	Integer	150	From the dimension table
zipcodeid	Zipcode id	Integer	500	From the dimension table
employerid	Employer id	Integer	1000	From the dimension table
occupationid	Occupation id	Integer	25	From the dimension table
dateid	Date id	Integer	45	From the dimension table
transactionid	Transactions id	Integer	1200	From the dimension table
committeeid	Committee id	Integer	60	From the dimension table
amount	Contribution amount by each contributor	Numeric(6,2)	-2838.81, 800.90, 600	Already available
Primary Key	Combination of stateid, cityid, zipcodeid, employerid, occupationid, dateid, transactionid, committeeid			

Foreign Keys	stateid REFERENCES id in State table cityid REFERENCES id in City table zipcodeid REFERENCES id in Zipcode table employerid REFERENCES id in Employer table occupationid REFERENCES id in Occupation table dateid REFERENCES id in Date table transactionid REFERENCES id in Transactions table committeid REFERENCES id in Committee_Name table
--------------	---

Please try use at least two types of input formats (SQL and text)

2.2 Hadoop Implementation

To be completed with project number 2

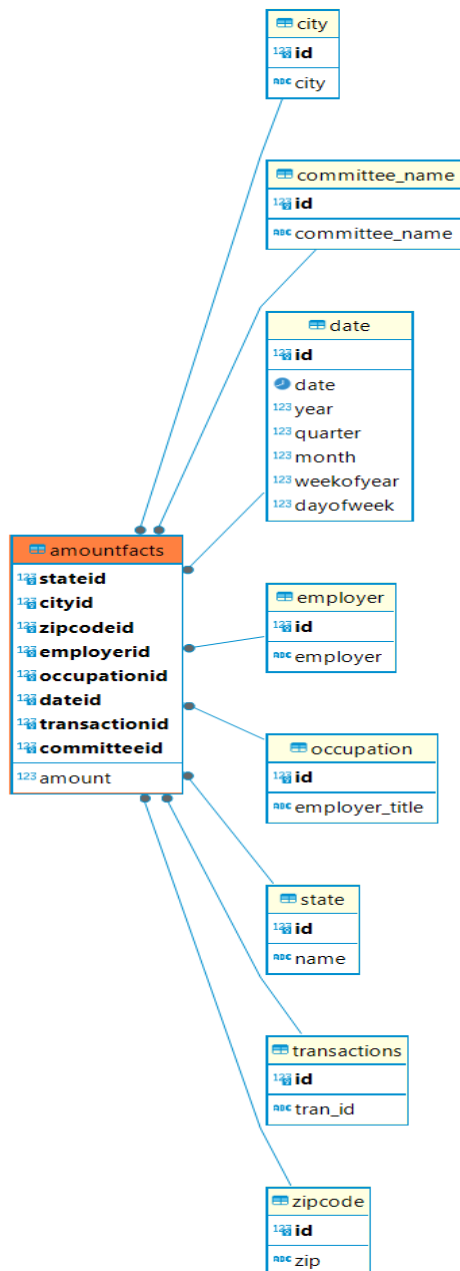
2.3 Reflective analysis of using a data warehouse vs Hadoop.

In project 1, please enter your reflective findings. Final version to be completed with project number 2.

My reflective findings during the design stage of this project include some discoveries of data types which PostgreSQL allows in its data management system. These values included Boolean types, Character Types (char, varchar, and text), Numeric types (such as integer and floating-point numbers), Temporal types (such as date, time, timestamp), and serial types. I found these findings interesting because I had the assumption that all data management systems used similar data types. And that was the case overall, but there were still some features that I was not aware of that PostgreSQL provided. One of these features included the data type serial.

Before completing this project, I was used to creating auto incrementing columns in tables through AUTO_INCREMENT or using the Identity modifier to increment a column. However, in PostgreSQL the serial data type is used to automatically increment a column value, so it preforms like a primary key. It is able to do this through the use of a regex expression that uses an auto generated sequence to increase the column's value sequentially.

I also found the use of indexing relevant in the design phase of data warehouses too. Since, I learned that indexing plays a crucial role in enhancing the performance and efficiency of database operations. Especially when one is designing tables in a data warehouse to handle data from large datasets. In this project I mainly used indexes on columns that served as primary keys or had unique constraints. This way PostgreSQL's data management system could have fast retrieval times for specific rows in these columns.



Indexing Values in tables:

Table Name:

city

Tablespace:

pg_default

Partition by:

Comment:

Object ID:

57346

Owner:

student

Extra Options:

Has Row-Level Security

Partitions

Columns	Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method	Tablespace	Comments
Constraints	city	city_city_idx_1	city	—	—	[]	—	—	296K	btree	pg_default	
Foreign Keys	city	city	—	[v]	[]	—	text_ops	—	—	—	—	
Indexes	id	city_pk_1	city	—	—	[v]	—	—	184K	btree	pg_default	
Dependencies	id	id	—	[v]	[]	—	int4_ops	—	—	—	—	

Table Name:	<input type="text" value="amountfacts"/>	Object ID:	<input type="text" value="57639"/>
Tablespace:	<input type="text" value="pg_default"/>	Owner:	student
<input type="checkbox"/> Has Row-Level Security <input type="checkbox"/> Partitions		Extra Options:	
Partition by:	<input type="text"/>		
Comment:	<input type="text"/>		

	Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method
Columns	dateid	amountfacts_dateid_idx_1	amountfacts	—	—	[]	—	—	604M	btree
Constraints	dateid	dateid	—	[v]	[]	—	int4_ops	—	—	—
Foreign Keys	dateid	dateid	—	—	—	—	—	—	—	—
Indexes	stateid, cityid, zipcodeid, employerid, occupationid	amountfacts_pk	amountfacts	—	—	[v]	—	—	1.8G	btree
	stateid	stateid	—	[v]	[]	—	int4_ops	—	—	—
	cityid	cityid	—	[v]	[]	—	int4_ops	—	—	—
	zipcodeid	zipcodeid	—	[v]	[]	—	int4_ops	—	—	—
	employerid	employerid	—	[v]	[]	—	int4_ops	—	—	—
	occupationid	occupationid	—	[v]	[]	—	int4_ops	—	—	—
	dateid	dateid	—	[v]	[]	—	int4_ops	—	—	—
	transactionid	transactionid	—	[v]	[]	—	int4_ops	—	—	—
	committeeid	committeeid	—	[v]	[]	—	int4_ops	—	—	—

Table Name:

committee_name

Object ID:

49223

Tablespace:

pg_default

Owner:

student

☐ Has Row-Level Security
 ☐ Partitions

Extra Options:

Partition by:

Comment:

	Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate
Columns	<div> <div>committee_name</div> <div>committee_name</div> </div>	<div> <div>committee_name_committee_name_idx</div> <div>committee_name</div> </div>	committee_name	—	—	[]	—	
Constraints	<div> <div>committee_name</div> <div>id</div> </div>	<div> <div>committee_name</div> <div>committee_name_pk</div> </div>	<div>—</div> <div>committee_name</div>	<div>[v]</div> <div>—</div>	<div>[]</div> <div>—</div>	<div>—</div> <div>[v]</div>	<div>text_ops</div> <div>—</div>	<div>—</div> <div>—</div>
Foreign Keys	id	id	—	[v]	[]	—	int4_ops	—
Indexes								
Dependencies								

Table Name:	date	Object ID:	57425
Tablespace:	pg_default	Owner:	student
<input type="checkbox"/> Has Row-Level Security <input type="checkbox"/> Partitions		Extra Options:	
Partition by:			
Comment:			

Columns	Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method
Constraints	id	date_pk_1	date	—	—	[v]	—	—	88K	btree
Foreign Keys	id	id	—	[v]	[]	—	int4_ops	—	—	—
Indexes	quarter, dayofweek	newtable_quarter_idx_1	date	—	—	[]	—	—	160K	btree
Dependencies	quarter	quarter	—	[v]	[]	—	int4_ops	—	—	—
	dayofweek	dayofweek	—	[v]	[]	—	int4_ops	—	—	—
References	weekofyear	weekofyear	—	[v]	[]	—	int4_ops	—	—	—
Partitions	month	month	—	[v]	[]	—	int4_ops	—	—	—

Table Name:

zipcode

Object ID:

57406

Tablespace:

pg_default

Owner:

student

☐ Has Row-Level Security
☐ Partitions

Extra Options:

Partition by:

Comment:

Columns

Constraints

Foreign Keys

Indexes

Dependencies

Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method	Tablespace
id	zipcode_pk_1	zipcode	—	—	[v]	—	—	208K	btree	pg_default
id	id	—	[v]	[]	—	int4_ops	—	—	—	—
zip	zipcode_zip_idx_1	zipcode	—	—	[]	—	—	288K	btree	pg_default
zip	zip	—	[v]	[]	—	text_ops	—	—	—	—

Table Name:	occupation	Object ID:	57370
Tablespace:	pg_default	Owner:	student
<input type="checkbox"/> Has Row-Level Security	<input type="checkbox"/> Partitions	Extra Options:	
Partition by:			
Comment:			

	Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method
Columns	employer_title	occupation_employer_title_idx_1	occupation	—	—	[]	—	—	1.4M	btree
Constraints	employer_title	employer_title	—	[v]	[]	—	text_ops	—	—	—
Foreign Keys	id	occupation_pk_1	occupation	—	—	[v]	—	—	720K	btree
Indexes	id	id	—	[v]	[]	—	int4_ops	—	—	—
Dependencies										

Table Name:

transactions

Object ID:

57394

Tablespace:

pg_default

Owner:

student

☐ Has Row-Level Security
 ☐ Partitions

Extra Options:

Partition by:

Comment:

Columns

Constraints

Foreign Keys

Indexes

Dependencies

Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method	Tablespace
id	transactions_pk_1	transactions	—	—	[v]	—	—	17M	btree	pg_default
id	id	—	[v]	[]	—	int4_ops	—	—	—	—
tran_id	transactions_tran_id_idx_1	transactions	—	—	[]	—	—	22M	btree	pg_default
tran_id	tran_id	—	[v]	[]	—	text_ops	—	—	—	—

Table Name:

state

Tablespace:

pg_default

☐ Has Row-Level Security
 ☐ Partitions

Partition by:

Comment:

Object ID:

57382

Owner:

student

Extra Options:

	Column	Index Name	Table	Ascending	Nullable	Unique	Operator Class	Predicate	Rel Size	Access Method	Tablespace
Columns	name	state_name_idx_1	state	—	—	[]	—	—	16K	btree	pg_default
Constraints	name	name	—	[v]	[]	—	text ops	—	—	—	—
Foreign Keys	id	state_pk_1	state	—	—	[v]	—	—	16K	btree	pg_default
Indexes	id	id	—	[v]	[]	—	int4 ops	—	—	—	—
Dependencies											

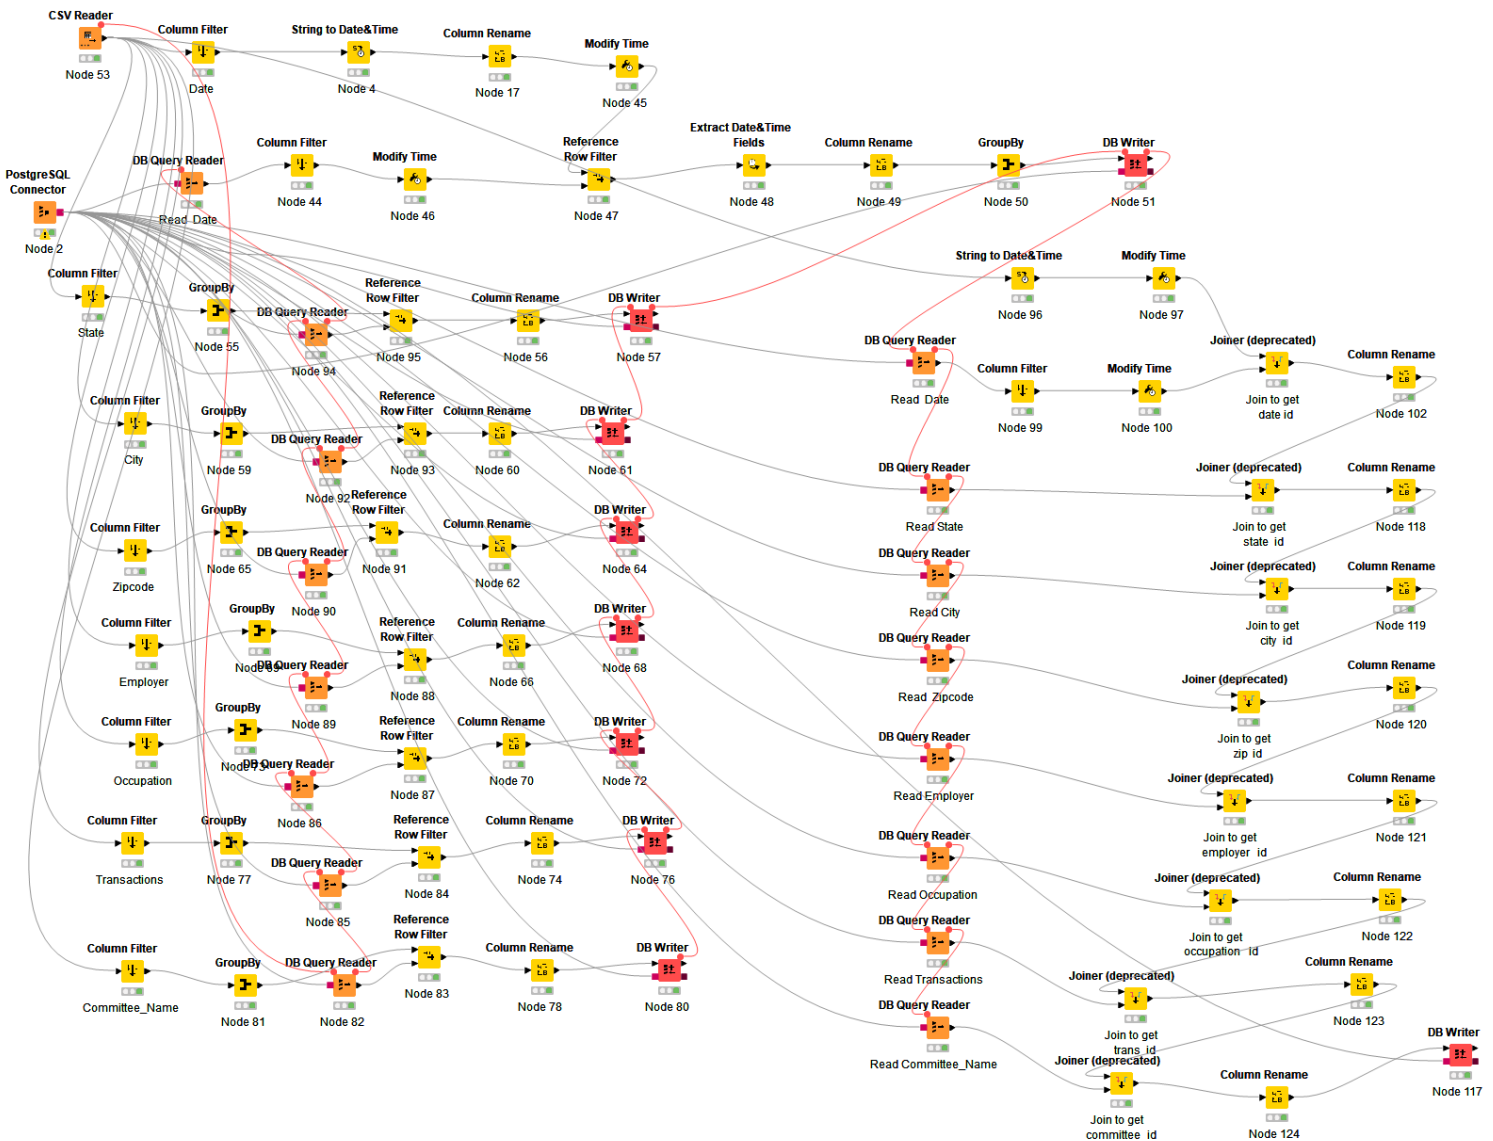
Overall Schema of Datawarehouse:

Properties	ER Diagram	Project1 Databases warehousepart1 Schemas warehouse							
Name:	warehouse	Namespace ID:	24577						
Comment:		Owner:	student						
Tables	Table Name	Object ID	Owner	Tablespace	Row Count Estimate	Has Row-Level Security	Partitions	Partition by	Extra Options
Foreign Tables	amountfacts	57,639	student	pg_default	3,805,618	[]	[]		
Views	city	57,346	student	pg_default	2,914	[]	[]		
Materialized Views	committee_name	49,223	student	pg_default	51	[]	[]		
Indexes	date	57,425	student	pg_default	1,520	[]	[]		
Functions	employer	57,358	student	pg_default	10,401	[]	[]		
Sequences	occupation	57,370	student	pg_default	5,716	[]	[]		
Data types	state	57,382	student	pg_default	0	[]	[]		
Aggregate functions	transactions	57,394	student	pg_default	93,050	[]	[]		
Permissions	zipcode	57,406	student	pg_default	3,104	[]	[]		
Source									

3. Data Preparation

3.1 Relational Data Warehouse Implementation ETL considerations

ETL Process Flow with description



As one could see this ETL process included many steps to draw data from the initial source (4 election contribution data files). Then transfer the data into the data

warehouse under warehousepart 1 project created from the student database under the schema warehouse in PostgreSQL.

Step 1: Extract the data from the input files

In the first step I used a CSV reader to access the data from each of the four comma delimited files stored on my computer of the four state's election contributions data.

The screenshot shows a CSV reader configuration window with the following sections:

- Input location:**
 - Read from: Local File System
 - Mode: ☒ File ☐ Files in folder
 - File: C:\Users\esbro\OneDrive\Desktop\DAAN825\First_Project\contributions_q1_2019_NJ.csv
- Reader options:**
 - Format:**
 - Autodetect format (checked)
 - Column delimiter: ,
 - Row delimiter: ☒ Line break ☐ Custom \r\n
 - Quote char: "
 - Quote escape char: "
 - Comment char: #
 - Has column header: ☒ (Has row ID: ☐)
 - Support short data rows: ☐ (Prepend file index to row ID: ☐)
- Preview:**
 - The suggested column types are based on the first 12000 rows only. See 'Advanced Settings' tab.
 - Table with 9 columns: Row ID, filing_id, linenum..., flag_or..., org_name, last_na..., first_na..., middle..., prefix.

Row ID	filing_id	linenum...	flag_or...	org_name	last_na...	first_na...	middle...	prefix
Row0	SA17A	SA17A	IND	?	Fuerst	Elizabeth	?	?
Row1	SA17A	SA17A	IND	?	Gillick	James	?	?
Row2	SA17A	SA17A	IND	?	Templo	Mariestella	?	?

Then I used a PostgreSQL connector to access and connect to the data in the data warehouse in PostgreSQL using the following credentials:

The screenshot shows a PostgreSQL connector configuration window with the following sections:

- Configuration:**
 - Database Dialect: PostgreSQL
 - Driver Name: PostgreSQL [ID: PostgreSQL]
- Location:**
 - Hostname: localhost
 - Database name: warehousepart1
- Authentication:**
 - ☐ Credentials
 - ☒ Username & password
 - Username: student
 - Password: *****
 - ☐ Kerberos

Step 2: Transform and Load data in the Dimension Tables

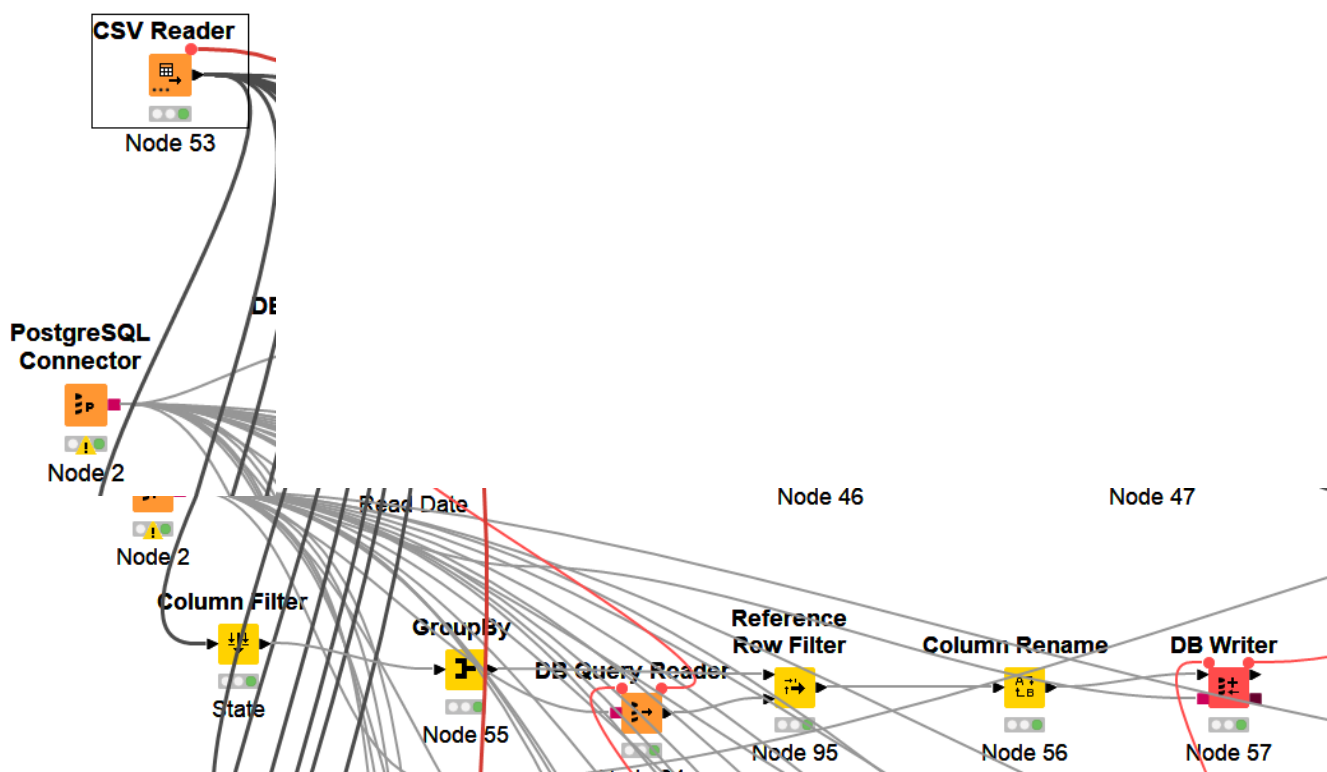
The I had to look at the dimensional data, for example State. We need to do the following for each dimension:

- Remove all other columns except for the dimension
- Keep only unique values for the dimension
- Save the dimensions in the PostgreSQL table

Note that you may get some errors if:

- You save the data file to a different path;
- Your Database runs under a different IP address

In Knime the corresponding operators are found in Manipulation. However you can search them by name. Below is the ETL flow for loading data in the “State” table.



Column Filter Flow Variables Job Manager Selection Memory Policy

☒ Manual Selection ☐ Wildcard/Regex Selection ☐ Type Selection

Exclude

Filter

- S fling_id
- S linenumbr
- S flag_origind
- S org_name
- S last_name
- S first_name
- S middle_name
- S prefix
- S suffix
- S address_one
- S address_two
- S city
- I zip
- S employer
- S occupation
- D amount
- S date
- D aggregate_amount
- S memo_code
- S memo_text
- S tran_id
- S back_ref_tran_id
- S back_ref_sched_name
- S prigen
- I cycle
- S fecid
- S committee_name

Include

Filter

- S state

Settings Description Flow Variables Job Manager Selection Memory Policy

Groups Manual Aggregation Pattern Based Aggregation Type Based Aggregation

Group settings

Available column(s)

Filter

Group column(s)

Filter

- S state

Change columns Flow Variables Job Manager Selection Memory Policy

Column Search

Filter Options: None

S state

state

Change: name StringValue

Settings Output Type Mapping Flow Variables Job Manager Selection Memory Policy

Table to write

Schema: warehouse Table: state

Batch Size: 1,000 ☒ Fail on error ☐ Append write status columns ☐ Disable DB Data output port ☐ Remove existing table

Select the columns to write (SET in SQL)

☒ Manual Selection ☐ Wildcard/Regex Selection ☐ Type Selection

Exclude

Filter

No column in this list

Include

Filter

- S name

I used a similar approach for the seven other dimensions that included city, zipcode, transactions, committee_name, employer, and occupation to load the data in the data warehouse. For date dimension I needed to follow a different approach due to the fact that dates are transient and every ETL will add new dates.

To execute the flow, right-click on the last node and click "Execute". You can view the data in DBeaver by right-clicking on the table and then click "View Data" and then "View Top 100 Rows".

Input Data with Write Status - 4:57 - DB Writer

File Edit Hilite Navigation View

Table "default" - Rows: 1 Spec - Column: 1 Properties Flow Variables

Row ID	\$ name
Row0	NJ

Table "default" - Rows: 767 Spec - Column: 1 Properties Flow Variables

Row ID	\$ city
Row0	ABSECON
Row1	ALLENDALE
Row2	ALLENHURST
Row3	ANDOVER
Row4	ASBURY PARK
Row5	ATCO
Row6	ATLANTIC HIGHLANDS
Row7	AUDUBON
Row8	AVENEL
Row9	AVON BY THE SEA
Row10	Aberdeen
Row11	Absecon
Row12	Allendale
Row13	Allentown

Row ID	I zip
Row0	0
Row1	7001
Row2	7002
Row3	7003
Row4	7004
Row5	7005
Row6	7006
Row7	7007
Row8	7008
Row9	7009
Row10	7010

Row ID	\$ employer_title
Row0	?
Row1	638 STEAMFITTER
Row2	ACCOUNT ANALYST
Row3	ACCOUNT EXECUTIVE
Row4	ACCOUNTANT
Row5	ACCOUNTING
Row6	ACTUARY
Row7	ADMIN ASSISTANT
Row8	ADMINISTRATION
Row9	ADMINISTRATIVE
Row10	ADMINISTRATIVE ASSISTANT
Row11	ADMINISTRATOR

Row ID	\$ employer
Row0	?
Row1	1919Investment Counsel
Row2	1975
Row3	21st Century Fox / Truex
Row4	221 Direct
Row5	24 EAST 73 LLC
Row6	39 New York Ave, LLC
Row7	7 Toy Drive
Row8	79 Hudson St LLC
Row9	A La Carte Premier Servers LLC
Row10	A.P.G. SECURITY, L.L.C.
Row11	A1A CLAIMS SVCS

Row ID	\$ tran_id
Row0	1000010
Row1	1000015
Row2	1000017
Row3	1000032
Row4	1000037
Row5	1000038
Row6	1000042
Row7	1000050
Row8	1000051

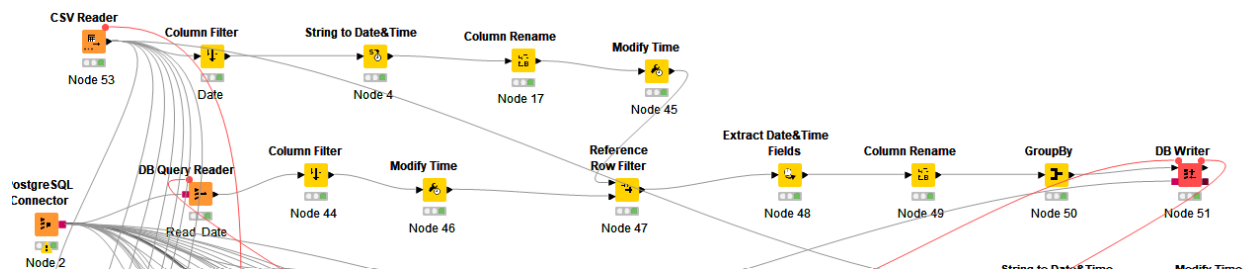
Table "default" - Rows: 16 Spec - Column: 1 Properties Flow Variables	
Row ID	committee_name
Row0	Amy For America
Row1	Bernie 2020
Row2	Beto For America
Row3	Cory 2020
Row4	Donald J. Trump For President, Inc.
Row5	Friends Of Andrew Yang
Row6	Friends Of John Delaney
Row7	Gillibrand 2020
Row8	Hickenlooper 2020
Row9	Inslee For America
Row10	Julian For The Future
Row11	Kamala Harris For The People
Row12	Marianne Williamson For President
Row13	Pete For America, Inc.
Row14	Tulsi Now
Row15	Warren For President, Inc.

Step 3: Load Data in the Date Dimension Table

To load data in the facts table, we need to:

- Load the date dimension into dimension table
- Change all the dimension names with dimensions ids.
- Remove data that is not needed

For this, we need to first read the data file and the dimension tables. The dimension tables need to be joined with the data from the data table so we can find the dimension id. The final workflow is shown below.



The most interesting operator here is the “Reference Row Filter” node that performs a “minus” operation on data, that is keeps only rows that exist in the first table but do not exist in the second table. That is necessary because we need to add only dimensions that do not already exist in the dimension table.

Another addition is the red line between the Excel Reader and the PostgreSQL Connector nodes on the left. This connection, noted in red line, will delay the execution

of the PostgreSQL connection till the successful completion of the Excel Reader node. It can be useful to get a new connection every time we change the input file. You can enable this by right-clicking on each node and selecting "Show Flow Variable Ports" and then connecting the ports.

Note: although not shown in this example, it is important that you avoid duplications in the dimension tables by using the "Reference Row filter" method.

Step 4: Load Data in the Facts Table

Once we have added the dates, we can load the data in the facts table. The flow at the very beginning of the ETL section of this project shows the flow for loading the election data in the warehouse schema data warehouse.

Again, the most interesting node is the "Joiner" node that we use to join the data with the dimension tables, on description/name to the id of the dimension. A second important operator is the String to date, that converts the string in the file to a date using a mask.

3.2 Hadoop Implementation

To be completed with project number 2

3.3 Reflective analysis of data preparation in relational data warehouse vs Hadoop.

In project 1, please enter your reflective findings. Final version to be completed with project number 2.

Using KNIME I was able to learn how to automate the ETL process through different nodes and functions to save time, reduce manual errors, and ensure that my data integration tasks were performed consistently and efficiently. I enjoyed that KNIME used a graphical interface to make the data extraction, transforming, and loading phases easier to understand, build, and test compared to other software applications I used to conduct this process.

Since the ETL process should be automated in most cases. And should be able to run at fixed times with limited human interaction. I believe I was able to accomplish this objective, since I only had to perform two steps to run the total ETL workflow in KNIME. These two steps included:

- 1: load each of the four respective new files into the csv node one at a time
- 2: execute the last DB Writer node on the right

4. Reporting System

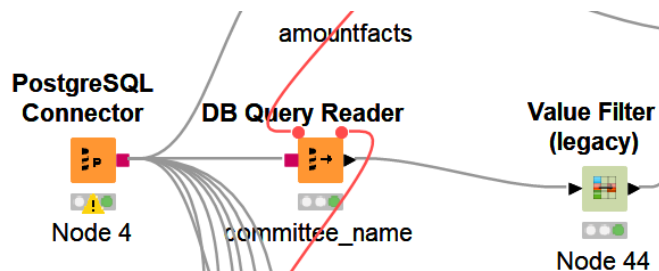
Provide screenshot and sample results with discussion on potential knowledge that was elicited.

4.1 Relational Data Warehouse Implementation

Retrieving and aggregating the data from the data warehouse

Our goal is to evaluate election contributions by state, political party, occupation, employer, month of the year, or day of year.

The first step is to create an interactive Knime workflow that will aggregate the data existing in the data warehouse. For this we will use the QuickForms “Value filter” that allows us to select the desired values from the database. For example, the flow with the three nodes below, establish a connection to the PostgreSQL server, read the data from the committee_name table, and then allows us to interactively select political party values.



The first node returns the id and committee name of the political party in the database table.

```
SQL Statement
1 select *
2 from warehouse."committee_name"
```

The second node allows us to define the column that is used for selection. As you can see in the screenshot below, I have chosen the “priority” column and I have locked it. I have also chosen to use all the values by default, see the “Include” panel.

Control Flow Variables Job Manager Selection Memory Policy

Label: Label

Description: Enter Description

Variable Name: new variable

Parameter Name: value-filter

Selection Type: Twinlist

Lock Column: ☒

Default Column: committee_name

Default Values:

Exclude

Filter

No values in this list

Include

Filter

>

>>

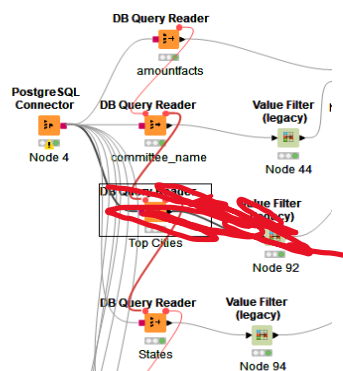
<

<<

- Amy For America
- Bernie 2020
- Beto For America
- Cory 2020
- Donald J. Trump For President, Inc.
- Friends Of Andrew Yang
- Friends Of John Delaney
- Gillibrand 2020
- Hickenlooper 2020
- Inslee For America
- Julian For The Future

The selection of the priority values is done by right clicking on the “Value Filter” node and then clicking “Execute and Open View”. Note that if the node is green you need to first reset the node to have the “Execute and Open View” available. The new window will give you the option to select the desired values. The new selected values will be applied when you click the button “Apply” and choose the “temporarily” option.

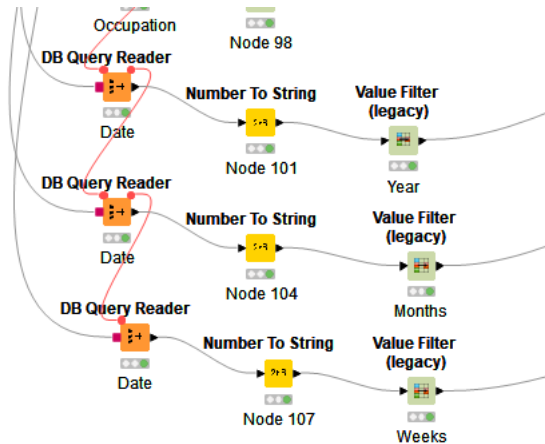
We repeat this process for two dimensions state and amountfacts tables, as shown in the figure below. The first node is reading from the facts table and it is not used yet. The following code is used to retrieve data from the database:



state Table: SELECT * FROM "state"

amountfacts Table: SELECT * FROM "amountfacts"

The data in the "Date" table needs to be processed some more due to the fact that the "Value Filter" node accepts only string columns/values while all the values in this table are numeric. To accomplish this we convert the Day of Month, day of week, and year to string. And then we feed them into the "value Filter" nodes.



The other three dimensions city, employer, and occupation I had use a slightly different retrieval method to get the proper data from these dimensions in the data warehouse.

```
Select Upper(c.city) as city, count, id
from
(select count(Upper(city)) as count, Upper(city) as city
from warehouse."city"
group by Upper(city)
having count(Upper(city))>=3
order by count(Upper(city)) desc) as countCity join warehouse."city" as c on
Upper(countCity.city)=Upper(c.city)
order by count desc;
```

For example for the "city" dimension I really only wanted to select city values that appeared in the table the most often. This is why I used a count function to count the number of occurrences of each city value in the city table. Then I only outputted values that occurred three times or more in the table. Which is basically outputting cities that contributed the most frequent during the 2020 elections in the first quarter.

I used this same concept for extracting data from the "occupation" table in the data warehouse, as well. I used count function to count the number of occurrences of each job title listed in the occupation table. Then I only outputted values that occurred seven

times or more in the table. Which is basically outputting job titles that contributed the most frequent during the 2020 elections in the first quarter.

SQL Statement

```

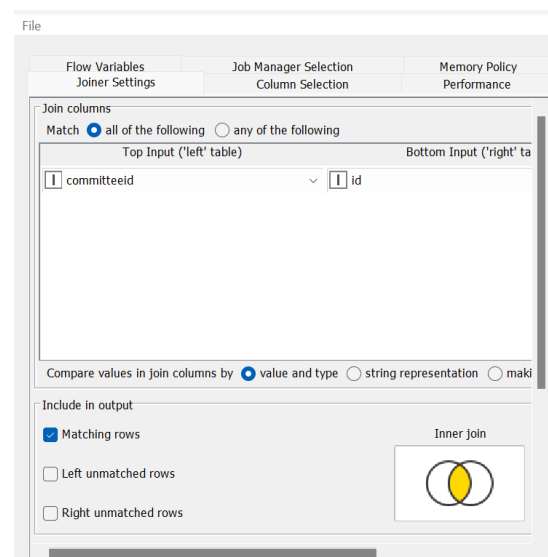
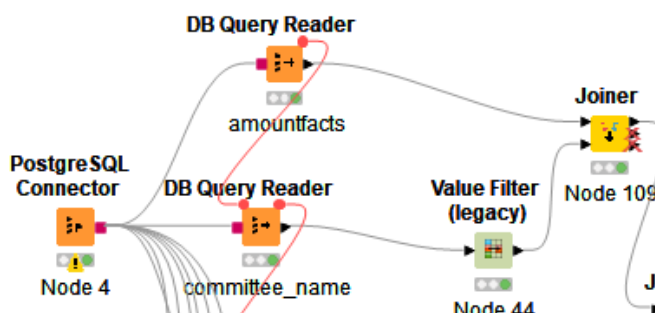
1  Select Upper(o.employer_title) as job_title, count, id
2  from
3  (select count(Upper(employer_title)) as count, Upper(employer_title) as Job_Title
4  from warehouse."occupation"
5  group by Upper(employer_title)
6  having count(Upper(employer_title))>=7
7  order by count(Upper(employer_title)) desc) as countOcc join warehouse."occupation" as o on
8  Upper(countOcc.Job_title)=Upper(o.employer_title)
9  order by count desc;

```

Lastly for the employer table I did not need to select certain values from this table. Since I wanted to get a full analysis of which employers contributed the most in 2020 elections. This did not mean the employers who contributed the most frequently.

employer Table: SELECT * FROM "employer"

Next, we will join the values in each dimension table with the fact table. For example, the flow below, joins the data in the facts table with the selected political parties. Joining is done on the pair "committeeid" from the facts table and the "id" field from the committee_name dimension table. We also have chosen to discard the id fields related to committee_name from the data set.



Top Input (left table)

Exclude

Filter

1 | committeeid

Enforce exclusion

Manual Selection Wildcard/Regex Selection Type Selection

Include

Filter

1 | stateid
1 | cityid
1 | zipcodeid
1 | employeoid
1 | occupationid
1 | datedid
1 | transactionid
1 | amount

Enforce inclusion

Bottom Input (right table)

Exclude

Filter

1 | id

Enforce exclusion

Manual Selection Wildcard/Regex Selection Type Selection

Include

Filter

8 | committee_name


Enforce inclusion

Duplicate column names


Do not execute

Append custom suffix (right)

We join the fact table with all the dimensions we have in a similar fashion. The result is shown in the figure below, together with a partial view of the resulting data.



The last three nodes aggregate the data by year, month of year, week of year, state, city, occupation, employer, and committee_name, and change the name of columns to more meaningful values.

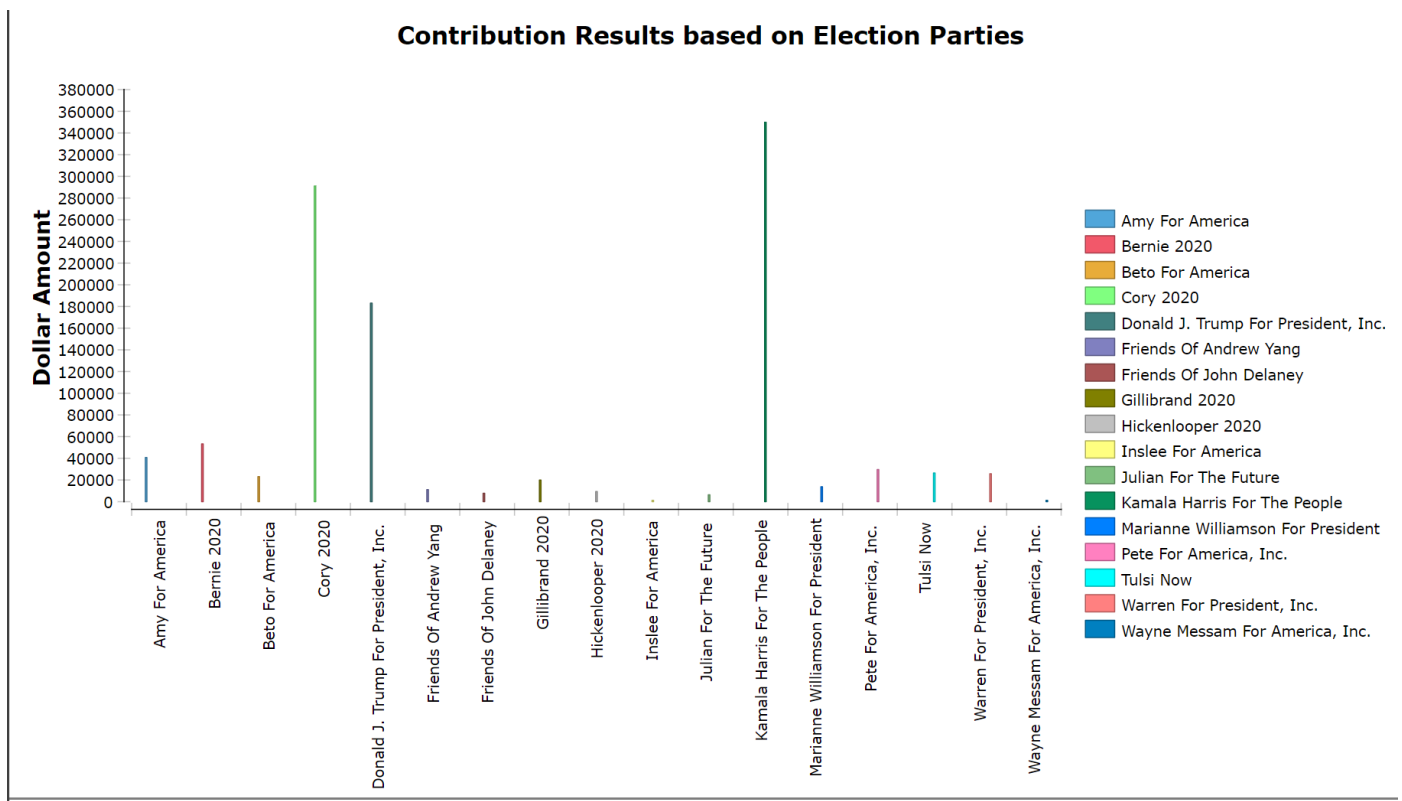


Creating a Report

Once we have added the report node and connected it to the data, we can customize it by clicking on the “Open Report” button on the menu.

Once the data was established to conduct the report building process, I began to create the necessary charts and tables needed to answer each of the six business questions. The first chart I began to create was used to answer the first business question that I established. Which was “Which campaigning party had the highest contributions in terms of dollars between the four states of California, New Jersey, Florida, and Arizona?”


This chart, I felt needed to be a bar chart since we wanted to get a visual understanding of what the top campaigning/political parties were based on the four states. And bar charts should be used when you are showing segments of information. Vertical bar charts are useful to compare different categorical or discrete variables (*Statistics of Canada, 2023*).



The second chart I began to create was used to answer the third business question that I established. Which was *“Which state had the highest contributions to the top two campaigning parties?”*

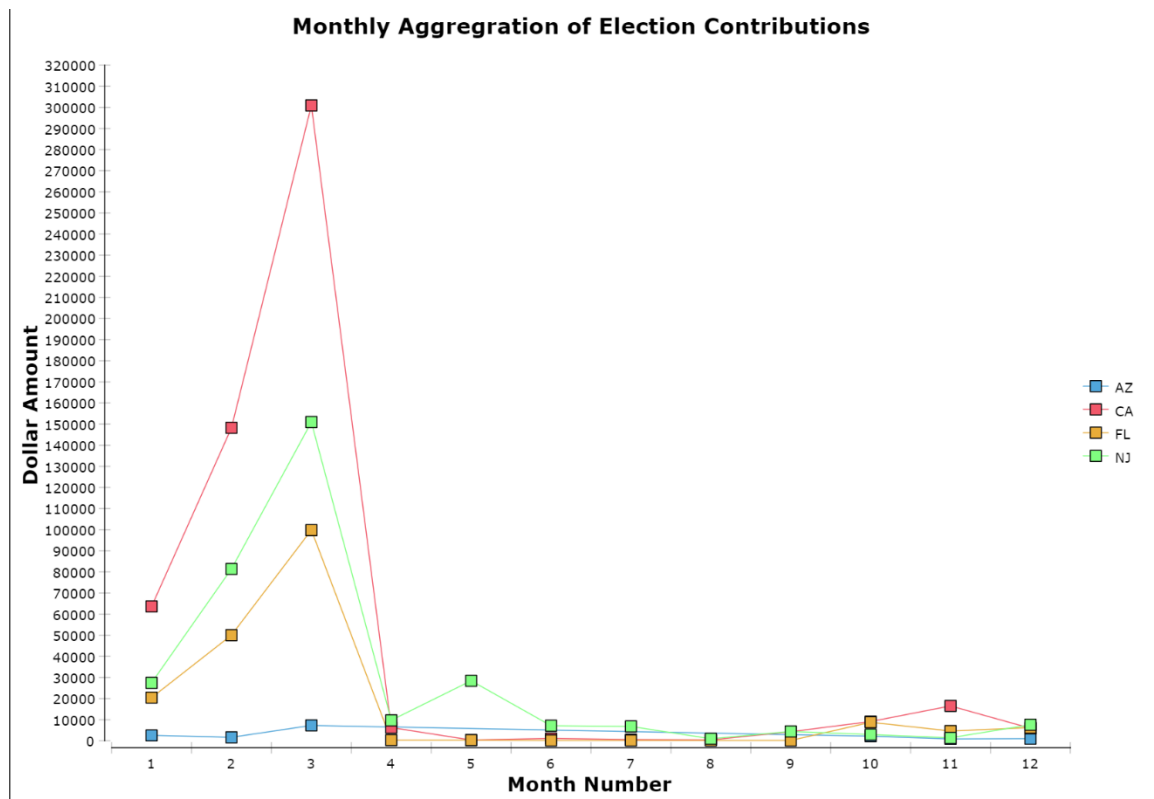
This chart, I felt, needed to be a pie chart since we wanted to get a visual understanding of which state contributed the most to the top political parties. And since the number of categorical variables that needed to be analyzed were less than five. I felt a pie chart could show this data point as a percentage of a whole. Instead of individual categorical aggregational values, which a bar chart would show.

State Contribution Distribution Between Top Two Election Parties



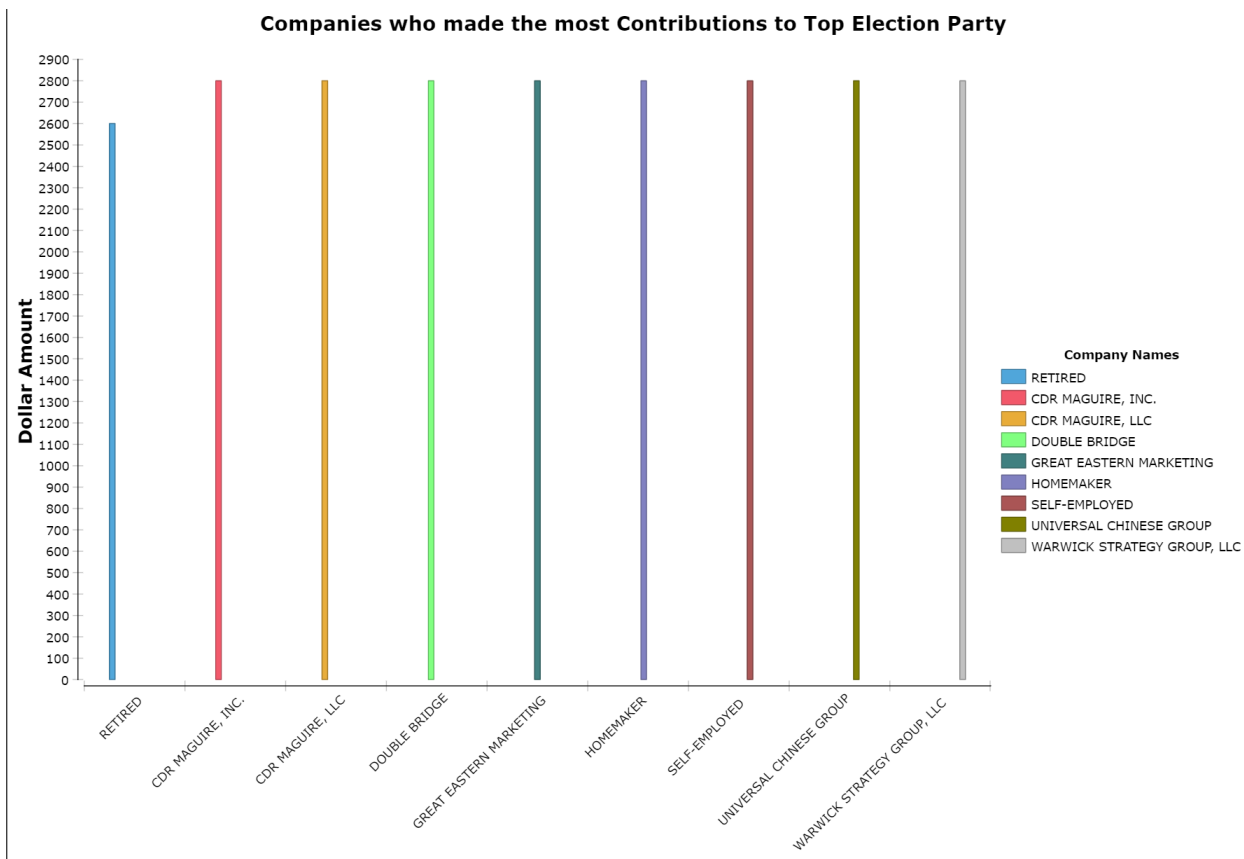
The third chart I began to create was used to answer the fourth business question that I established. Which was *“Display contributions made in California, New Jersey, Florida, and Arizona aggregated monthly in Q1 during the year of 2019.”*

This chart, I felt, needed to be a line graph since we wanted to get a visual understanding of how contributions by individuals changed over time in the first quarter of elections in 2019. And line graphs are best used when showing how something changes over time (KidsZone, 2023).



The fourth chart I began to create was used to answer the fifth business question that I established. Which was *“Which company made the most contributions in terms of dollars to the top campaigning party?”*

To answer this question, I felt I could choose from numerous different chart types. Though from further analysis I felt a bar chart would best depict the visual needed to answer this question. Since there was a great number of companies used to create the dataset for this report. And I felt a pie chart would not show the percentage of contributions from a given company compared to other companies in the dataset best. Because there would be too many companies to analyze and compare against with the use of slices in a pie chart. Also, a line graph would not suffice because we were not analyzing the change in values over a particular time period.



The fifth and final chart I began to create was used to answer the sixth business question that I established. Which was *“What was the overall demographic makeup of the contributors of each of the top campaigning parties?”*

For this question I had to think twice about what I actually wanted to show to the audience of the report. And I came to the conclusion that not a graph or plot would best describe the demographic makeup of the contributors of the top campaigning parties. There would be too much information needed to be displayed to describe a given supporter of a given political party. So, I decided to use a table to show all the different type political party supporters for each of top campaigning parties. The information I decided to show the table would include the different job professions each political party had supporting their campaign in the 2020 elections.

Political Party	Occupation	Dollar Amount
Donald J. Trump For President, Inc.	ACCOUNTANT	300
	ADMINISTRATOR	60
	ANALYST	20
	ARCHITECT	73
	ARTIST	1108
	ATTORNEY	62
	BANKER	9
	BOOKKEEPER	25
	BUSINESS OWNER	27
	CEO	208
	CONSTRUCTION	74
	CONSULTANT	680
	CONTRACTOR	35
	DENTIST	114
	DESIGNER	75
	DIRECTOR	83
	DRIVER	73
	ENGINEER	93
	ENTREPRENEUR	60
	EXECUTIVE	67
	HOMEMAKER	194
	INFORMATION REQUESTED	219
	INVESTOR	777
	LAWYER	146
	MANAGER	107
	MANAGING DIRECTOR	85
	MARKETING	27
	NOT EMPLOYED	35
	OWNER	57
	PARALEGAL	0
	PHARMACIST	35
	PHYSICIAN	54
	PRESIDENT	343
	PRINCIPAL	2800
	PROFESSOR	17
	PROJECT MANAGER	56
	REAL ESTATE	63
	REAL ESTATE BROKER	75
	REALTOR	30
	REGISTERED NURSE	56
	RETIRED	62
	SALES	143
	SELF EMPLOYED	169
	SELF-EMPLOYED	156
	SOFTWARE ENGINEER	500
	STUDENT	63
	SURGEON	82
	TEACHER	70
	TECHNICIAN	19
	TRUCK DRIVER	52
	VICE PRESIDENT	432
Total Contribution Amount		183100

<i>Total Contribution Amount</i>	Political Party	Occupation	183100	Dollar Amount
Kamala Harris For The People		ACCOUNTANT	750	
		ADMINISTRATOR	100	
		ANALYST	132	
		ARCHITECT	916	
		ARTIST	144	
		ATTORNEY	1161	
		BANKER	1000	
		BUSINESS OWNER	2800	
		CEO	1903	
		CONSULTANT	462	
		CONTRACTOR	250	
		DENTIST	2800	
		DESIGNER	189	
		ENGINEER	527	
		EXECUTIVE	1877	
		EXECUTIVE DIRECTOR	1183	
		HOMEMAKER	1520	
		INVESTOR	2800	
		LAWYER	1433	
		MANAGER	749	
		MANAGING DIRECTOR	250	
		MARKETING	1400	
		NOT EMPLOYED	250	
		OWNER	1857	
		PHYSICIAN	1284	
		PRESIDENT	1062	
		PROFESSOR	237	
		REAL ESTATE	2000	
		REAL ESTATE BROKER	500	
		REALTOR	375	
		REGISTERED NURSE	250	
		RETIRED	297	
		SALES	465	
		SELF EMPLOYED	625	
		SOFTWARE ENGINEER	184	
		STUDENT	161	
		VICE PRESIDENT	750	
<i>Total Contribution Amount</i>			349765	

<i>Total Contribution Amount</i>	Political Party	Occupation		Dollar Amount
Cory 2020		ADMINISTRATOR	860	
		ARCHITECT	1266	
		ARTIST	168	
		ATTORNEY	960	
		BANKER	33	
		BOOKKEEPER	200	
		CEO	837	
		CONSULTANT	1208	
		CONTRACTOR	1000	
		DESIGNER	25	
		DIRECTOR	2700	
		EXECUTIVE	1294	
		HOMEMAKER	950	
		INFORMATION REQUESTED	883	
		INVESTOR	657	
		MANAGER	54	
		MANAGING DIRECTOR	453	
		NOT EMPLOYED	86	
		OWNER	773	
		PHYSICIAN	287	
		PRESIDENT	985	
		PRINCIPAL	250	
		PROFESSOR	102	
		PROJECT MANAGER	1000	
		REAL ESTATE	782	
		REAL ESTATE BROKER	500	
		REALTOR	1000	
		RETIRED	1118	
		SALES	594	
		SELF-EMPLOYED	1050	
		STUDENT	2700	
		SURGEON	2800	
		TEACHER	500	
		VICE PRESIDENT	25	
<i>Total Contribution Amount</i>			291009	

Lastly, to answer the second business question of “*Which campaigning party had the lowest contributions in terms of dollars amongst the four states of California, New Jersey, Florida, and Arizona?*”, I decided that I did not need to create a sixth and final chart. Mainly because my first bar chart of campaign contributions showed this finding.

4.2 Hadoop Implementation

To be completed with project number 2

4.3 Reflective analysis of result in relational data warehouse vs Hadoop.

In project 1, please enter your reflective findings. Final version to be completed with project number 2.

Conclusions

Overall conclusions of the project. In project 2, add a reflective analysis of the advantages and disadvantages of the two implementations.

When analyzing and answering the business question stated earlier in this project assignment, I was able to discover some interesting facts. One point I was able to discover was the campaigning party who had the highest contributions in terms of dollars between the four states was Kamala Harris. Second was Cory Booker and third was Donald J. Trump. Kamala Harris had raised around \$350,000, Cory Booker raised around \$291,000, and Donald J. Trump raised around \$183,000 from the four states. The lowest campaigning party amongst the four states was a tie between Jay Inslee and Wayne Messam. Both raised only \$1,250 dollars amongst the four states.

Then the state who had the highest contribution value amongst the top two political candidates of Donald J. Trump and Kamala Harris was California. With a total of \$320,903.96 raised in the state. Florida came in second, New Jersey third, and Arizona last.

Next, in terms of monthly aggregations of contributions from each of the four states, I was able to see there was a spike of contributions from January to March. Contributions went from around \$64,000 to \$301,000 during this time period. However, after March contributions amongst the four states came crashing down to around \$10,000 on

average for the remainder of the year.

Then in terms of companies who contributed the most to Kamala Harris campaign were people who were retired, or worked for the following companies: CDR Maguire, Double Bridge, Great Eastern Marking, Homemaker, people self-employed, Universal Chinese Group, or Warwick Strategy Group. All these companies and/or people on average contributed around \$2,800 to her campaign.

Lastly in terms of demographic makeup of the top campaigning candidates. Of which people who contributed the most were attorneys, accountants, artists, principals, real estate agents, executive directors, dentists, CEOs, business owners, bankers, people who worked in marketing, physicians, architects, surgeons, or students.

Works Cited

Government of Canada, Statistics Canada. “5 Data Visualization 5.2 Bar Chart.” *5.2 Bar Chart*, 2 Sept. 2021, www150.statcan.gc.ca/n1/edu/power-pouvoir/ch9/bargraph-diagrammeabarres/5214818-eng.htm.

“Line Graph.” *Kids Graphing Page - Line Graph - NCES Kids’ Zone*, 14 Oct. 2023, nces.ed.gov/nceskids/graphing/classic/line.asp#:~:text=Line%20graphs%20can%20be%20used,for%20what%20is%20being%20measured.

DBeaver documentation (2023) *DBeaver*. Available at: <https://dbeaver.com/docs/dbeaver/> (Accessed: 10 October 2023).

KNIME documentation (2023) *KNIME*. Available at: <https://docs.knime.com/> (Accessed: 2 October 2023).

“The Collaborative Data Science Platform | Mode.” *Mode.com*, 2 Oct. 2023, mode.com/. Accessed 2 Oct. 2023.